

MAKING VALID INFERENCE WITH DECISION TREE

GEORGE EKOW QUAYE

Master's Program in Statistics

APPROVED:

Xiaogang Su, Ph.D., Chair

Suneel B. Chatla, Ph.D.

Wen-Yee Lee, Ph.D.

Stephen L. Crites, Jr., Ph.D.
Dean of the Graduate School

©Copyright

George Ekow Quaye

2021

To my

*FATHER Carlos Quaye, MOTIVATOR Dr. Jacob Setorglo and my ADVISOR Dr.
Xiaogang Su*

with love

MAKING VALID INFERENCE WITH DECISION TREE

by

GEORGE EKOW QUAYE

THESIS

Presented to the Faculty of the Graduate School of

The University of Texas at El Paso

in Partial Fulfillment

of the Requirements

for the Degree of

MASTER OF SCIENCE

Department of Mathematical Sciences

THE UNIVERSITY OF TEXAS AT EL PASO

May 2021

Acknowledgements

First and foremost I am extremely grateful to my supervisor, Dr. Xiaogang Su for his invaluable advice, continuous support, and patience during my study. His immense knowledge and plentiful experience have encouraged me in all the time of my academic research and daily life. My gratitude extends to the faculty of Mathematical Science for the funding opportunity to undertake my studies at the Department of Mathematical Sciences, the University of Texas at El Paso. Additionally, I would like to express gratitude to Dr. Suneel B Chatla and Dr. Wen-Yee Lee for their treasured support which was really influential in shaping my thesis fundings and critiquing my results. I also thank Dr. Jacob Setorglo of the University of Cape Coast and Mr. Carlos Quaye for their mentorship, without their tremendous understanding and encouragement in the past few years, it would be impossible for me to complete my study. Finally to my family and friends for their encouragement and support all through my studies.

Abstract

Hypothesis testing and Confidence Interval (CI) estimates are key statistics in predicting future values in data analysis. Most often, CI estimates are directly obtained from the summary statistics of a particular statistical methodology output. However, when it comes to the summary of decision tree outputs, these CI estimates are not directly obtained. So a naïve way of making node-level inference is to construct a $(1 - \alpha) \times 100\%$ confidence interval for a node mean \bar{y}_t using the relation: $\bar{y}_t \pm z_{1-\alpha/2} \frac{s_t}{\sqrt{n_t}}$, where \bar{y}_t is the node mean and s_t is the standard deviation estimates from the decision tree summary. Nevertheless, these sets of intervals tend to be over-optimistic owing to the very adaptive nature of tree modeling, in other words, they are too narrow to have the desired coverage. This challenge with CI in tree summary stands as one of the most common requests from the users of decision trees that are however rarely fulfilled in practice. In this research, we make a strong effort to nail out the source of over-optimistic and correct it accordingly. We began by treating this issue with an existing method known as the Bootstrap Calibration (BC) on the α . Statistically, this BC method is also plagued with overfitted estimates. We then resorted to our approach (Bootstrap Bias Correction), an approach that seeks to correct a downwards biasedness in the s_t estimates to obtained bias-corrected SD estimates (s_t''). Now ,the node mean \bar{y}_t , the node sample size n_t , a fixed α value together with the BBC estimate s_t'' was then used to obtain a more accurate CI intervals for \bar{y}_t through the relation: $\bar{y}_t \pm z_{1-\alpha/2} s_t^{(')} / \sqrt{n_t}$. The CI estimates from the proposed method (BBC) were empirically assessed and illustrated through simulation studies and validated via real data exploration.

Table of Contents

	Page
Acknowledgements	v
Abstract	vi
Table of Contents	vii
List of Tables	ix
List of Figures	x
Chapter	
1 Introduction	1
1.1 Background	1
1.2 Problem Statement	2
1.3 Overview of Thesis	2
2 Literature Review	4
2.1 Decision Trees	4
2.2 Types of Decision Trees	6
2.3 Extension of Decision Trees and Recent development	7
2.4 Problem with Inference	8
3 Methodology	10
3.1 Motivation	10
3.1.1 Bias Correction Approach	11
3.1.2 Bootstrapping	12
3.2 Exiting Methods	13
3.2.1 Bootstrap Calibration on α	13
3.3 Proposed Method	16
3.3.1 Bootstrap Bias Correction on s_t	17
4 Simulations Study	20

4.1	Bootstrap Calibration (BC)	22
4.2	Proposed Method	24
4.2.1	Bias Correction on SD	24
4.2.2	Empirical Coverage	30
4.2.3	Percentage Coverage Estimates	32
5	Real Data Analysis	33
5.1	Real Data Exploration	33
5.1.1	Data Source and Preparation	33
5.1.2	Obtaining the best Tree from the data set	36
5.1.3	Estimates of s_t and s_t^c from the best tree	37
5.1.4	Empirical Coverage from the best tree	37
6	Discussion and Conclusion	39
6.1	Significance of the Result	39
6.2	Recommendation	40
Appendix		
	Appendix	41
	Vita	59

List of Tables

4.1	Mean values, Influence of tree modeling on inference on node averages \bar{y}_t and node SD s_t for $t \in \tilde{T}$ computed with training data \mathcal{D} and test data \mathcal{D}' . . .	21
4.2	Table of result for biased corrected SD computed with training data \mathcal{D} and test data \mathcal{D}'	24
4.3	Table of result for biased corrected SD computed with training data \mathcal{D} and test data \mathcal{D}'	26
4.4	Table of result for biased corrected SD computed with training data \mathcal{D} and test data \mathcal{D}'	28
4.5	Confident intervals using the s_t , s_t'' and s_t' from data generated by Model A	31
4.6	Confident intervals using the s_t , s_t'' and s_t' from data generated by Model B	31
4.7	Confident intervals using the s_t , s_t'' and s_t' from data generated by Model C	31
4.8	Coverage Probabilities by BBC	32
5.1	1987 Baseball Salary Data for Hitters	35
5.2	Table values of s_t and s_t'' using the baseball salary data.	37
5.3	Confidence estimates of the mean salary via BBC: $\bar{y}_t \pm z_{1-\alpha/2}s_t''/\sqrt{n_t}$. . .	37

List of Figures

2.1	Decision tree presenting response to direct mailing. Sourced from <i>Decision trees in Data mining and knowledge discovery handbook</i> (pages 165–192), by Rokach, L. and Maimon, O. (2005) Springer.	6
3.1	Influence of tree modeling on inference: (a) node averages \bar{y}_t and (b) node SD s_t for $t \in \tilde{T}$ computed with training data \mathcal{D} and test data \mathcal{D}' . The green reference line is $y = x$	11
4.1	Coverage plot of the calibrated alpha against population alpha.	23
4.2	Illustration of bias correction on SD s_t through model A. The green reference line is $y = x$	25
4.3	Illustration of bias correction on SD s_t through model B. The green reference line is $y = x$	27
4.4	Illustration of bias correction on SD s_t through model C. The green reference line is $y = x$	29
5.1	Plot of the best tree via pruning and cross validation.	36
5.2	Plot of the Confident Estimates of the Mean Salary	38

Chapter 1

Introduction

1.1 Background

Decision trees are a well known statistical tool in machine learning and statistics for predictive analysis (e.g. classification and regression) (Lakshminarayanan, 2016). According to Lakshminarayanan (2016) learning a decision tree from training data involves training the tree structure \mathcal{T} , estimating the leaf node parameters Ω , and predicting a label within each leaf node. Well-known decision tree algorithms include CART (Breiman et al., 1984) and C4.5 (Quinlan, 1993). Decision trees come with several advantages in practical applications such as the ability to be well-suited for datasets with mixed attribute types (e.g. binary, categorical, real-valued attributes) and interpretability (at least on simple problems) (Lakshminarayanan, 2016). Although decision trees are powerful and advantageous in some way, they are prone to over-fitting (over-optimism) and require several studies to limit their complexity in order to minimize their generalization predictive error. Obtaining a decision tree model \mathcal{T} consists of three components: a method of splitting data, a method of determining the best tree model, and a method of summarizing the terminal node. With regards to the third component, it is common that only the node size and the mean response, i.e., $\{n_t, \bar{y}_t\}$, are included as node summary at each terminal node $t \in \tilde{\mathcal{T}}$, where $\tilde{\mathcal{T}}$ denotes the set of all the terminal nodes of \mathcal{T} . Note that \bar{y}_t amounts to the proportion of 1's in the case of classification trees, on which basis the majority rule can be used to produce the 0/1 summary. One difficulty with this summarizing method is that it does not allow for statistical inference such as a confidence interval for the true node mean μ_t . Though inference stands as one of the most common requests from the users of decision trees, this

issue has rarely been fulfilled in practice.

1.2 Problem Statement

A naïve way of making node-level inference is to construct a $(1 - \alpha) \times 100\%$ confidence interval (CI)

$$\bar{y}_t \pm z_{1-\alpha/2} \frac{s_t}{\sqrt{n_t}} \quad (1.1)$$

for each terminal node $t \in \tilde{\mathcal{T}}$, where $z_{1-\alpha/2}$ is $(1 - \alpha/2)$ -th percentile of the standard normal $\mathcal{N}(0, 1)$ distribution and s_t denotes the standard deviation (SD) of responses in node t and n_t for the node size. Nevertheless, these sets of intervals are over-optimistic owing to the very adaptive nature of tree modeling, in other words, they are too narrow to have the desired coverage .

1.3 Overview of Thesis

Chapter 1 talks briefly about the decision trees, their derivation process, and the over-optimism problem when used for prediction purposes, which has been a long-standing issue for scientists. Chapter 2 further elaborates on decision trees including their history, types, extensions, and recent developments such as interaction trees and oblique decision trees and their problem with statistical inference or prediction. In subsequent chapters, we describe in detail the source of the prediction problem and available methods in treating that. Specifically, chapter 3 discusses source of overoptimism, available approaches in treating this issue and our proposed method or algorithm which seeks to outperform the existing method. A simulation study on decision tree modeling is carried out in chapter 4, where we further elaborate the optimism problem given the three models, investigate the performance of our proposed method and compare it with competitive approaches. In chapter 5 we perform an illustrative example using real data on the obtained accurate

proposed method to illustrate the use of our method in a practical setting. Chapter 6 which is the final chapter summarizes the thesis and discusses possible future research directions.

Chapter 2

Literature Review

A tree-based method (or recursive partitioning) divides data recursively to attain multiple mutually exclusive sub-groups. Tree-based methods are very effective in handling multifaceted data and acquiring acknowledgment as a sound technique for addressing data complexity, which renders them attractive in different application fields. The proposal of Classification and Regression Tree (CART) (Breiman et al., 1984) made the tree models more popular and widely accepted in applications and the current norm of tree modeling.

2.1 Decision Trees

A decision tree is a graphical representation of specific decision situations that are used when complex branching occurs in a structured decision (Njoku, 2019). While in data mining a decision tree is a predictive model which can be used to represent both classifiers and regression models, in operations research decision trees refer to a hierarchical model of decisions and their consequences (Maimon and Rokach, 2014). The implementation of decision trees originated from decision theory and statistics. One of the best and most applied supervised learning algorithm in predictive modeling is decision trees which also works in connection with ensemble methods for more accurate results. Decision trees are general purpose prediction and classification mechanisms that were among the first statistical algorithms to be implemented in electronic form during the adoption of digital circuitry to electronic computations in the latter decades of the 20th century (de Ville, 2013). According to Hu et al. (2019), Decision trees are one of the leading forms of interpretable models and despite several attempts over the last several decades to improve the opti-

mality of decision tree algorithms, the CART (Breiman et al., 1984) and C4.5 (Quinlan, 1993) decision tree algorithms (and other greedy tree-growing variants) have remained as dominant methods in practice. Obtaining a decision tree model according to the CART (Breiman et al., 1984) convention involves growing a large initial tree, a pruning algorithm for reducing the tree size, and a validation method for determining the best tree. Tree methods are also an excellent tool for grouping. Once a final tree structure is obtained, the groups are naturally induced by its terminal nodes.

The decision tree consists of three types of nodes that are a root node that has no incoming edge, an internal or test node that has exactly one incoming edge, and outgoing edges and leaves (terminal or decision nodes). According to a given function of the input attributes values, each internal node splits an instance space into two or more sub-spaces on a decision tree.

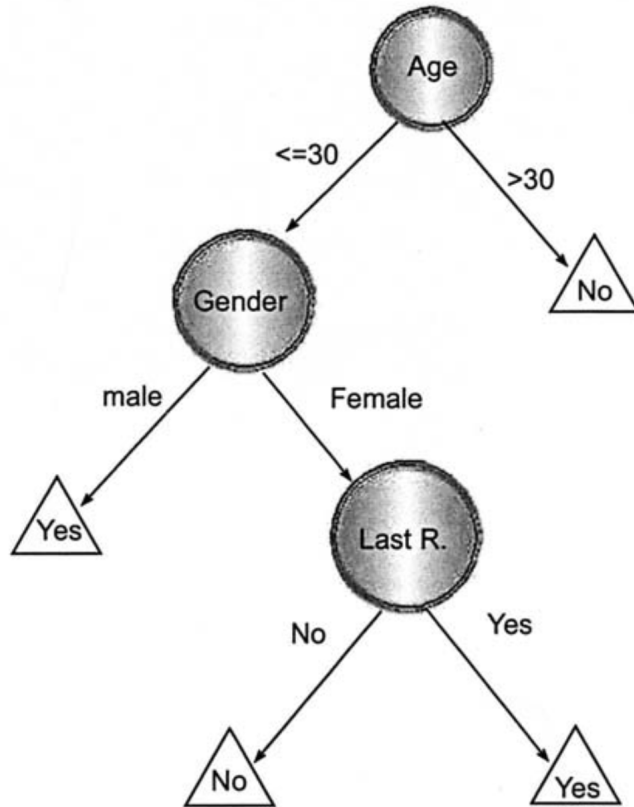


Figure 2.1: Decision tree presenting response to direct mailing. Sourced from *Decision trees in Data mining and knowledge discovery handbook* (pages 165–192), by Rokach, L. and Maimon, O. (2005) Springer.

Figure 2.1 presents a decision tree that shows whether or not a potential customer will respond to a direct mailing. Given Figure 2.1, one can predict the response of a potential customer by sorting it down the tree and understand the behavioral characteristics of the entire potential customers population regarding direct mailing (Rokach and Maimon, 2005).

2.2 Types of Decision Trees

There are basically two types of trees;

- Regression trees when the outcome or response variable is continuous, for example

determining the price of a newly manufactured product by considering the various inputs and constraints.

- Classification trees (also known as decision trees) when the outcome is binary or categorical. A classical example is a toss of a coin which has only two outcomes whether a head or tail.

From a statistical perspective, regression as a whole is a broader concept that incorporates the classification problem as a special case.

2.3 Extension of Decision Trees and Recent development

Decision trees as originally coined from tree models have undergone several recent developments. Oblique decision trees according to Murthy and Salzberg (1995) produce polygonal (polyhedral) partitionings of the attribute space, while conventional axis-parallel trees produce partitionings in the form of hyper-rectangles that are parallel to the feature axes. A general-purpose data structure for addressing the behaviors of recursive programs that interact with their surroundings is known as interaction trees (ITrees) (Xia et al., 2019). Interaction trees (ITrees) were employed by Su et al. (2008) in their article, where ITrees was to optimize a subgroup analysis in comparative studies. Specifically, the IT method recursively partitions the data into two subsets that show the greatest interaction with the treatment, which results in a number of objectively defined subgroups (Su et al., 2008).

Decision trees have as well generated several extensions. Noticeable amongst them are Multivariate Adaptive Regression Splines (MARS), Hierarchical Mixture Model (HMM), and the Ensemble Methods (EM). Multivariate Adaptive Regression Splines (MARS) (Friedman, 1991) is a new method presented for flexible regression modeling of high dimensional data and this procedure is motivated by the recursive partitioning approach to regression and shares its attractive properties. While CART does piecewise constant modeling, MARS

fits piecewise linear models. Also, Jordan and Jacobs (1994) presented a tree-structured architecture for supervised learning and the statistical model underlying the architecture is a hierarchical mixture model in which both the mixture coefficients and the mixture components are generalized linear models (GLIM's). Ensemble methods are machine learning technique that produces one optimal predictive model by combining usually hundreds or thousands base learners. By considering one decision tree and guessing to make the right decision at each split, ensemble methods make provision to take into consideration a sample of decision trees, evaluate which characteristics to employ or questions to enumerate at each split, and then make a final prediction as a result of the aggregated results of the sampled decision trees. Noticeable ensemble methods used in obtaining an optimal outcome in decision trees are Boosting, Bagging and Random Forests(RF). Bagging is a method used to improve on unstable estimators or classifiers in a learning model. Specifically by generating multiple versions of the classifier and using these to get an aggregated classifier to obtain the new model (Breiman, 1996). Boosting (Freund et al., 1996) serves as a tool to significantly minimize the error of any learning algorithm that consistently generates classifiers that are better than guessing randomly. Random Forest (RF) is a combination of trees such that each tree depends on independently random sampled vector values of the same distribution Breiman (2001).

2.4 Problem with Inference

One major difficulty in summarizing decision trees is that it does not allow for statistical inferences such as confidence interval or hypothesis testing for the true node mean μ_t . A naive approach in making a node inference is therefore to construct a $(1 - \alpha) \times 100\%$ confidence interval (CI) or perform a hypothesis test using the two stochastics components \bar{y}_t and s_t from the tree summary. This approach, however, leads to an over-optimism of the estimates or too narrow confidence intervals. An accurate or prudent way of making a valid inference within terminal nodes of decision trees has been a long-standing challenge

for statisticians. Among few works that have been done to overcome this challenge, Loh et al. (2018) proposed using the bootstrap calibration (BC) approach (Loh, 1987, 1991) to tune the confidence level. Particularly a best α' is sought such that $(1 - \alpha')$ intervals of the same form as (1.1) has the $(1 - \alpha)$ coverage for each terminal node.

Chapter 3

Methodology

3.1 Motivation

We began by first making efforts to understand the influence of tree modeling as a data-adaptive approach on statistical inference and identify the source of overoptimism. The naïve confidence interval (CI) estimates in (1.1) involves two stochastic components \bar{y}_t and s_t , the sample mean and sample standard deviation computed with observations in the training data \mathcal{D} that fall into terminal node t . In the following, we designed a study to investigate the performance of these two components in estimating the true node mean and node SD.

We generate training data \mathcal{D} of size $n = 500$ from one nonlinear model (Friedman, 1991):

$$y = -6 + 0.1 \exp(4x_1) + 4 \exp\{20(x_2 - 0.5)\} + 3x_3 + 2x_4 + x_5 + \varepsilon \quad (3.1)$$

with $\varepsilon \sim \mathcal{N}(0, 1)$ and the \mathbf{X}_i 's are generated independently from random uniform[0,1] distribution of size $n = 500$. A best-sized tree \mathcal{T} is then constructed via pruning and cross-validation with the 1-SE rule (Breiman et al., 1984). For each terminal node, \bar{y}_t and s_t are computed and recorded. Then we generate another independent test data set \mathcal{D}' of size $n' = 10,000$. We send \mathcal{D}' down to tree \mathcal{T} and recompute the node mean and SD (\bar{y}'_t, s'_t).




Figure 3.1: Influence of tree modeling on inference: (a) node averages \bar{y}_t and (b) node SD s_t for $t \in \tilde{T}$ computed with training data \mathcal{D} and test data \mathcal{D}' . The green reference line is $y = x$.

Figure 3.1 plots \bar{y}_t vs. \bar{y}'_t in Panel (a) and s_t vs. s'_t in Panel (b). It can be seen that \bar{y}_t and \bar{y}'_t match well with each other, indicating the node averages can be used for prediction purposes. Nevertheless, s_t from the training data \mathcal{D} are generally smaller than s'_t from the test data \mathcal{D}' , highlighting a systematic downward bias. Part of the reason accounting for the underestimated s_t is that data is split with greedy search by minimizing the within-node impurity or variation when building up the tree model. Using s_t directly would inevitably lead to inflated Type I error rates that hold accountable for the over-optimism. This observation in the s_t bias distribution motivates us to correct the bias in the SD estimator s_t .

3.1.1 Bias Correction Approach

There are several techniques and approaches in dealing with bias correction. According to Jiao and Han (2017) some general approaches to bias correction are the bootstrap, the jackknife, and the Taylor series. The jackknife uses a subsampling approach where

the biases of estimators with different sample sizes are made to cancel each other while the bootstrap bias correction on the other hand uses the plug-in rule to estimate the bias. Taylor series is used to create an estimate (guess) of what a function or population parameter looks like through a derivative at a single point. However, the Taylor series according to (Jiao and Han, 2017) is a less versatile method compared to the bootstrap and jackknife owing to its applicability to functions with specific global differentiability conditions. Given the natural grouping at each terminal node of decision trees, a small disturbance to the data set leads to different node membership observations. With this underlying behavior of trees and these methods discussed, bootstrapping (Tibshirani and Efron, 1993) was chosen for our study.

3.1.2 Bootstrapping

The bootstrapping approach was introduced by Efron (1979) among others with the motive of determining the variations in statistics when a theoretical variance is either unknown or not estimable and also correcting some forms of biasedness. The bootstrap methodology or concept is to simulate from an empirical distribution of a given data by means of resampling with replacement to obtain an approximation of the sampling distribution of the statistics. The bootstrap methodology estimates the sampling distribution of a given function Θ by recalculating its overall bootstrap samples B_1, \dots, B_n to obtain a set of bootstrapped statistics $\Theta_1, \Theta_2, \dots, \Theta_B$, where B is the number of bootstrap samples for a given set of data and a statistic of importance Θ . The required approximate estimate can then be estimated by using $\Theta_1, \dots, \Theta_B$. The bias is obtained by taking the difference between the average of the bootstrapped statistics $\Theta_1, \dots, \Theta_B$ and the given statistic Θ . Let Θ_b represent the set of $\Theta_1, \Theta_2, \dots, \Theta_B$, then;

$$Bias = \left(\frac{1}{B} \sum_{b=1}^B \Theta_b \right) - \Theta \quad (3.2)$$

Bias-corrected statistic is Θ^c now derived as:

$$\begin{aligned}\Theta^c &= \Theta - \text{Bias} \\ &= 2\Theta - \left(\frac{1}{B} \sum_{b=1}^B \Theta_b\right)\end{aligned}$$

Algorithm 1 shows the iteration for a simple learning algorithm for a bootstrap-bias correction problem.

<pre> input : data $\mathbf{X}_i \in \mathbb{R}^p$, $size = n$, Θ output: Acquire bias-corrected estimate Θ^c 1 begin 2 for $1 \rightarrow b$ to B do 3 Derive $\mathbf{X}_1, \dots, \mathbf{X}_B$ by resampling \mathbf{X}_i with replacement. Calculate $\Theta_1, \dots, \Theta_B$ 4 Estimate Θ^c, with $\Theta_1, \dots, \Theta_B$ 5 end 6 Acquire the bias-corrected estimate: 7 $\Theta^c = 2\Theta - (\frac{1}{B} \sum_{b=1}^B \Theta_b)$, where $\Theta_b \in (\Theta_1, \dots, \Theta_B)$ 8 end </pre>
--

Algorithm 1: Bootstrap Method

3.2 Exiting Methods

3.2.1 Bootstrap Calibration on α

Given that scientists and investigators are often interested in making simultaneous inferences across all terminal nodes of \mathcal{T} , e.g., construct simultaneous confidence intervals for the node means. To deal with the multiplicity issue and over-optimism of tree modeling, Loh et al. (2018) proposed using the bootstrap calibration (BC) approach (Loh, 1987, 1991)

to tune the confidence level. Specifically, a constant $0 < \alpha' < 1$ is sought such that $(1 - \alpha')$ intervals of the same form as (1.1) has the $(1 - \alpha)$ coverage for all terminal nodes. The BC algorithm in Loh et al. (2018) was originally designed for tree-structured subgroup analysis (see, e.g., Su et al., 2009), where differential treatment effects are of the major concern. Applying the same procedure to ordinary regression trees leads to Algorithm 2.


```

input : data  $\mathcal{D} = \{(\mathbf{x}_i, y_i) \in \mathbb{R}^p \times \mathbb{R}\}_{i=1}^n$  and  $0 < \alpha_1 < \dots < \alpha_K < \alpha < 1$  for
        a given  $\alpha$ .

output: calibrated confidence coefficient  $(1 - \alpha')$ .

1 initialize  $B$  – # bootstrap samples;
2 begin
3   Set coverage  $\gamma_k = 0$  for  $k = 1, \dots, K$  ;
4   for  $b \leftarrow 1$  to  $B$  do
5     draw a bootstrap sample  $\mathcal{D}_b$  ;
6     construct a best-sized tree  $\mathcal{T}_b$  from  $\mathcal{D}_b$  via pruning and cross validation;
7     summarize terminal nodes of  $\mathcal{T}_b$  as  $\{(n_t, \bar{y}_t, s_t) : t \in \tilde{\mathcal{T}}_b\}$  ;
8     send  $\mathcal{D}$  down to  $\mathcal{T}_b$  and recompute the mean response  $\bar{y}'_t$  on basis of  $\mathcal{D}$  ;
9     set counters  $c_k = 0$  for  $k = 1, \dots, K$  ;
10    for  $t \in \tilde{\mathcal{T}}_b$  do
11      for  $k \leftarrow 1$  to  $K$  do
12        construct  $(1 - \alpha_k) \times 100\%$  CI in node  $t$ :
13         $(L_{tk}, U_{tk}) \leftarrow \bar{y}_t \pm z_{1-\alpha_k/2} \frac{s_t}{\sqrt{n_t}}$  , if  $\bar{y}'_t \in (L_{tk}, U_{tk})$ ,  $\forall t$  then
14           $c_k := c_k + 1$ ;
15        end
16      end
17      update  $\gamma_k := \gamma_k + \frac{c_k}{|\tilde{\mathcal{T}}_b|}$  and average  $\gamma_k := \gamma_k/B$  for  $k = 1, \dots, K$ ;
18    end
19    find  $k^* \leftarrow$  smallest  $k$  such that  $\gamma_k < (1 - \alpha)$ , implying that  $\gamma_{k^*-1} \geq (1 - \alpha)$ 
        and obtain  $\alpha'$  via linear interpolation
        
$$\alpha' = \alpha_{k^*-1} + \frac{(1 - \alpha) - \gamma_{k^*}}{\gamma_{k^*-1} - \gamma_{k^*}} (\alpha_{k^*} - \alpha_{k^*-1}).$$

20 end

```

Algorithm 2: Bootstrap Calibration (BC)

Now, given \mathcal{D} as the data set and \mathcal{D}_b to be a random bootstrap sample from \mathcal{D} . Let \mathcal{T}_b be the set of all terminal nodes obtained from \mathcal{D}_b . For any $t \in \tilde{\mathcal{T}}_b$, let \bar{y}'_t be the node mean. We construct $(1 - \alpha')$ confidence intervals for \bar{y}'_t such that, the coverage probability of $(\mathcal{D}, \mathcal{D}_b, t, \alpha')$, averaged over the terminal nodes in the tree constructed from \mathcal{D}_b has expected value $(1 - \alpha)$.

Let γ_k denote the set coverage that encompasses a terminal node of \mathcal{T} of each observation in dataset \mathcal{D} , such that $P(\bar{y}'_t \in CI_{\gamma_k}) = (1 - \alpha) \forall t \in \tilde{\mathcal{T}}_b$, initializing $\gamma_k = 0$ for $k = 1, \dots, K$. We take B bootstrap samples $\{\mathcal{D}_b : b = 1, \dots, B\}$. For each bootstrap sample \mathcal{D}_b , a best-sized tree \mathcal{T}_b is constructed via pruning and cross validation to obtain the estimates of \mathcal{T}_b such as $\{(n_t, \bar{y}_t, s_t) : t \in \tilde{\mathcal{T}}_b\}$ for all terminal nodes of \mathcal{T}_b . Sending \mathcal{D} down to \mathcal{T}_b and recomputing $\{\bar{y}'_t : t \in \tilde{\mathcal{T}}_b\}$ for all terminal nodes of \mathcal{T}_b . A counter coverage $c_k = 0$ for $k = 1, \dots, K$ for each terminal node is set and we construct $(1 - \alpha_k) \times 100\%$ CI in node t for \bar{y}_t based on set coverage α_k in Algorithm 2. If $\bar{y}'_t \in (L_{tk}, U_{tk}) \in (\bar{y}'_t \pm z_{1-\alpha_k/2} \frac{s_t}{\sqrt{n_t}})$, then we update the set coverage $c_k = c_k + 1$ for every mean node. We further update the γ_k by averaging c_k over the total number of trees and adding it to the initial coverage set γ_k as outlined in line 18 of Algorithm 2. Subsequently we average γ_k over the number of bootstrap samples to obtain the smallest $K(k^*)$, such that $\gamma_k < (1 - \alpha)$, implying that $\gamma_{k^*-1} \geq (1 - \alpha)$. Now with the smallest k and its corresponding γ_k and α_k values our desired bootstrapped calibrated α' is obtained via linear interpolation.

3.3 Proposed Method

A statistically obtained prediction interval relies on the data \mathcal{D} . Making an reliable $(1 - \alpha) \times 100\%$ confidence prediction base on the data relies on the stochastic estimates \bar{y}_t and s_t from the summary of decision tree output. However, directly constructing this confidence interval with these estimates most especially the s_t turns to be over-optimistic as shown in Figure 3.1. Therefore we propose a method that keeps the α constant and corrects the downward biasedness in s_t through bootstrapping to obtain a more honest unbiased

estimate.

3.3.1 Bootstrap Bias Correction on s_t

In the bootstrap bias correction approach discussed, the estimates from bootstrap samples are compared to the original estimate and the averaged difference furnishes an estimator of the bias. However, there is one major obstacle with tree modeling. Trees are unstable in the sense that a small perturbation to the data often results in a substantially different tree model structure at the end. As a result, the tree models obtained with bootstrap samples are different from each other and from the final tree model constructed with the original sample. Hence to tackle this problem, we note that every tree model forms a natural grouping of the entire data. With two tree structures, observations in a node from one tree can be distributed into different nodes of the other tree. Utilizing this property, we put forward this feasible bootstrap bias correction procedure for the underestimated standard deviation s_t as outlined in Algorithm 3.

```

input : data  $\mathcal{D} = \{(\mathbf{x}_i, y_i) \in \mathbb{R}^p \times \mathbb{R}\}_{i=1}^n$ .
output: A tree model  $\mathcal{T}$  with bias corrected SD  $s_t$  for each  $t \in \tilde{\mathcal{T}}$ .

1 initialize  $B$  – # bootstrap samples;
2 begin
3   construct a best-sized tree  $\mathcal{T}$  from  $\mathcal{D}$  via pruning and cross validation;
4   obtain node membership vector  $\mathbf{m}_0 \in \mathbb{R}^n$  for all observations in  $\mathcal{D}$ 
     w.r.t.  $\mathcal{T}$ ;
5   compute  $s_t$  for each  $t \in \tilde{\mathcal{T}}$  based on  $\mathcal{D}$ ;
6   set bias  $b_t = 0$  for  $t \in \tilde{\mathcal{T}}$ ;
7   for  $b \leftarrow 1$  to  $B$  do
8     draw a bootstrap sample  $\mathcal{D}_b$ ;
9     construct a best-sized tree  $\mathcal{T}_b$  from  $\mathcal{D}_b$  via pruning and cross
       validation;
10    compute SD  $\{s_{bt'} : t' \in \tilde{\mathcal{T}}_b\}$  based on  $\mathcal{D}_b$ ;
11    send  $\mathcal{D}$  down to  $\mathcal{T}_b$  and recompute SD  $\{s_{0t'} : t' \in \tilde{\mathcal{T}}_b\}$  on basis of  $\mathcal{D}$ ;
12    compute bias  $b_{bt'} = s_{0t'} - s_{bt'}$  for  $t' \in \tilde{\mathcal{T}}_b$ ;
13    obtain node membership vector  $\mathbf{m}_b \in \mathbb{R}^n$  for all observations in  $\mathcal{D}$ 
       w.r.t.  $\mathcal{T}_b$ ;
14    form two-way contingency table  $\{m_{tt'} : t \in \tilde{\mathcal{T}} \text{ and } t' \in \tilde{\mathcal{T}}_b\}$  with  $\mathbf{m}_0$ 
       and  $\mathbf{m}_b$ ;
15    compute row proportions  $p_{tt'} = m_{tt'}/m_t$ ;
16    for  $t \in \tilde{\mathcal{T}}$  do
17      | update  $b_t := b_t + \sum_{t' \in \tilde{\mathcal{T}}_b} p_{tt'} b_{bt'}$ ;
18    end
19  end
20  average bias  $b_t := b_t/B$  for  $t \in \tilde{\mathcal{T}}$ ;
21  bias correction  $s_t'' := s_t + b_t$  for  $t \in \tilde{\mathcal{T}}$ .
22 end

```

Algorithm 3: Bias correction for SD in tree modeling.

Let $\mathbf{m}_0 \in \mathbb{R}^n$ denote the node membership vector that assigns a terminal node of \mathcal{T} to each observation in \mathcal{D} . We take B bootstrap samples $\{\mathcal{D}_b : b = 1, \dots, B\}$. For each bootstrap sample \mathcal{D}_b , a best-sized tree \mathcal{T}_b is constructed via pruning and cross validation to obtain the standard deviation estimates $\{s_{bt'} : t' \in \tilde{\mathcal{T}}_b\}$ for all terminal nodes of \mathcal{T}_b . Also we began by initially setting the bias $b_t = 0$ for $t \in \tilde{\mathcal{T}}$. Sending \mathcal{D} down to \mathcal{T}_b and recomputing $\{s_{0t'} : t' \in \tilde{\mathcal{T}}_b\}$ based on \mathcal{D} yield bias estimates $\{b_{bt'} = s_{0t'} - s_{bt'} : t' \in \tilde{\mathcal{T}}_b\}$ for each terminal node of \mathcal{T}_b . Our goal, however, is to obtain bias estimates b_t for each s_t in $\{s_t : t \in \tilde{\mathcal{T}}\}$. We do so with a weighted average of $b_{bt'}$ by looking at how observations in $t \in \tilde{\mathcal{T}}$ are distributed over $\tilde{\mathcal{T}}_b$. To proceed, let $\mathbf{m}_b \in \mathbb{R}^n$ denote the node membership vector that assigns a terminal node of \mathcal{T}_b to each observation in \mathcal{D} . The two categorical vectors \mathbf{m}_0 and \mathbf{m}_b form a $|\tilde{\mathcal{T}}| \times |\tilde{\mathcal{T}}_b|$ two-way contingency table with counts $\{m_{tt'} : t \in \tilde{\mathcal{T}} \text{ and } t' \in \tilde{\mathcal{T}}_b\}$. Let $p_{tt'} = m_{tt'}/m_{t\cdot}$ be the row marginal proportions, where $m_{t\cdot} = \sum_{t'} m_{tt'}$ is the t -th row total. Then an estimate of the bias from the b th bootstrap sample \mathcal{D}_b is given by $\sum_{t' \in \tilde{\mathcal{T}}_b} p_{tt'} b_{bt'}$. Averaging over B bootstrap samples leads to a bias estimate for s_t and bias correction on s_t can be made accordingly. Put together, the bias-corrected SD s_t'' is given by

$$s_t'' := s_t + \frac{1}{B} \sum_{i=1}^B \sum_{t' \in \tilde{\mathcal{T}}_b} p_{tt'} (s_{0t'} - s_{bt'}). \quad (3.3)$$

With valid s_t'' values from equation 3.3, one convenient way of summarizing terminal nodes could be simply to state $\{n_t, \bar{y}_t, s_t\}$ and leave the subsequent inferences (individual or simultaneous) to the users. Individual CI's, formula (3.4) with bias corrected s_t'' would suit.

$$\bar{y}_t \pm z_{1-\alpha/2} \frac{s_t''}{\sqrt{n_t}} \quad (3.4)$$

For simultaneous inferences, Bonferroni, FDA (false discovery rate), and other types of adjustment can be applicable. Especially if $|\tilde{\mathcal{T}}|$ is small or moderate, Bonferroni is appealing.

Chapter 4

Simulations Study

This section outlines simulation studies performed to analyze and investigate the source of overoptimism and how well the proposed methodology works in elevating this problem. To better understand the influences of adaptive methods on each stochastic components \bar{y}_t and s_t , we generate training data \mathcal{D} of size $n = 500$ and test data set \mathcal{D}' of size $n' = 10,000$ from two nonlinear models in (Friedman, 1991) and one true tree model;

$$y = -6 + 0.1 \exp(4x_1) + 4 \exp\{20(x_2 - 0.5)\} + 3x_3 + 2x_4 + x_5 + \varepsilon \quad (4.1)$$

$$y = 10 \sin(\pi x_1 x_2) + 20(x_3 - 0.5)^2 + 10x_4 + x_5 + \varepsilon \quad (4.2)$$

$$y = 2 + 2 \times \text{sign}(x_1 \leq 0.5) \times \text{sign}(x_2 \leq 0.5) + \varepsilon \quad (4.3)$$

with $\varepsilon \sim \mathcal{N}(0, 1)$. and the \mathbf{X}_i 's are generated independently from random uniform[0,1] distribution of size $n = 500$.

Model (4.1) has a nonlinear additive structure on the first two variables and a linear term on the last three variables (Friedman, 1991). Model (4.2) has a nonlinear additive with the first two variables having a multiplicative parabolic interaction term, the third variable with a quadratic relation, and the final two variables with a linear dependence (Friedman, 1991) and the true tree model is given by model (4.3). For simplicity, model equations 4.1, 4.2, and 4.3 would be denoted as models A, B, and C respectively.

With each given model, a best-sized tree \mathcal{T} is then constructed via pruning and cross-validation with the 1-SE rule (Breiman et al., 1984). For each terminal node, \bar{y}_t and s_t are extracted. Then we generate another independent test data set \mathcal{D}' of size $n' = 10,000$. Send \mathcal{D}' down to tree \mathcal{T} and recompute the node mean and SD (\bar{y}'_t, s'_t). Mean values of \bar{y}_t ,

s_t , \bar{y}'_t and s'_t are computed as indicated in Table 4.1

Table 4.1: Mean values, Influence of tree modeling on inference on node averages \bar{y}_t and node SD s_t for $t \in \tilde{T}$ computed with training data \mathcal{D} and test data \mathcal{D}' .

Tree	\bar{y}_t	s_t	\bar{y}'_t	s'_t
1	7.05608	1.328307	7.089057	1.50636
2	6.813485	1.335768	6.785658	1.471886
3	6.559332	1.226844	6.597213	1.43782
4	6.72372	1.220396	6.691766	1.437119
5	6.990564	1.362192	6.88289	1.480317
\vdots	\vdots	\vdots	\vdots	\vdots
96	6.818623	1.232034	6.895311	1.457672
97	6.435047	1.206669	6.425234	1.467036
98	7.006038	1.285327	6.986894	1.503891
99	6.904065	1.218327	6.846934	1.422443
100	6.878547	1.219453	6.752506	1.379429

It is observed from Table 4.1 given by equation model 4.1 that, on average \bar{y}_t and \bar{y}'_t has a correspondence to each other that is to say they match well in figures. However, the SDs exhibits some variations in figures, i.e, s'_t from the test data is much greater than that of the training data as indicated in Table 4.1. These average s_t values clearly support the vast variation in the standard deviations resulting in the downwards biasedness as depicted in figure 3.1. This observation further strengthened our motivation for the study.

4.1 Bootstrap Calibration (BC)

With the existing methodology BC and its outlined algorithm. We simulated data from the three models and use Algorithm 2 to obtain the coverage probabilities. Even though the population α is rarely known in practice, by simulation we exploit the luxury of data availability in the simulation setting and obtain an estimate of α by generating large test samples evaluation of the bootstrap calibration approach to aid our evaluation of the bootstrap calibration approach.

For each model configuration, we started by generating a training data \mathcal{D} of size $n = 500$ and also test data set \mathcal{D}' of size $n' = 500$. A set of $\alpha \in [1 : 0.005]$ are chosen. With the training data \mathcal{D} , a best-sized tree \mathcal{T} is then constructed via pruning and cross-validation with the 1-SE rule (Breiman et al., 1984) and the estimates $\{(n_t, \bar{y}_t, s_t) : t \in \tilde{\mathcal{T}}\}$ for all terminal nodes of \mathcal{T} are recorded. Send \mathcal{D}' down to \mathcal{T} and recomputing $\{\bar{y}'_t : t \in \tilde{\mathcal{T}}\}$ for all terminal nodes of \mathcal{T} . Now are we construct $(1 - \alpha_k) \times 100\%$ CI in node t for \bar{y}_t based on set coverages α . If $\bar{y}'_t \in (L_{tk}, U_{tk}) \in (\bar{y}'_t \pm z_{1-\alpha/2} \frac{s_t}{\sqrt{n_t}})$, then we record the α 's for which $\bar{y}'_t \in (L_{tk}, U_{tk}) \in (\bar{y}'_t \pm z_{1-\alpha/2} \frac{s_t}{\sqrt{n_t}})$ for every mean node as the population coverage probabilities. Similarly, we take B bootstrap samples $\{\mathcal{D}_b : b = 1, \dots, B\}$, such that for each bootstrap sample \mathcal{D}_b , a best-sized tree \mathcal{T}_b is constructed via pruning and cross validation to obtain the estimates of \mathcal{T}_b as $\{(n_t, \bar{y}_t, s_t) : t \in \tilde{\mathcal{T}}_b\}$ for all terminal nodes of \mathcal{T}_b .

Sending \mathcal{D} down to \mathcal{T}_b and recomputing $\{\bar{y}'_t : t \in \tilde{\mathcal{T}}_b\}$ for all terminal nodes of \mathcal{T}_b , a $(1 - \alpha_k) \times 100\%$ CI in node t for \bar{y}'_t based on set coverages α is then constructed. If $\bar{y}'_t \in (L_{tk}, U_{tk}) \in (\bar{y}'_t \pm z_{1-\alpha/2} \frac{s_t}{\sqrt{n_t}})$, then we record the α 's for which $\bar{y}'_t \in (L_{tk}, U_{tk}) \in (\bar{y}'_t \pm z_{1-\alpha/2} \frac{s_t}{\sqrt{n_t}})$ for every mean node as the bootstrap coverage probabilities. A plot of comparison for the population and bootstrap coverage probailities obtained in Figure 4.1.




Figure 4.1: Coverage plot of the calibrated alpha against population alpha.

From Figure 4.1, the dash line serves as a reference line with $(1 - \alpha) = 0.95$. Given the plot, our goal is that the calibrated (bootstrapped) alpha (α') should mimic that of the unknown population alpha (α) and their intersection points with the reference line should be close to each other. The graphical output indicates that model A and model B do not show our desire result, thus the bootstrap calibrated alpha is too liberal compared to the population alpha which is not good. The third model which is the true tree model shows good results, this is because its a true classification model. Hence the existing methodology BC becomes too radical to tackle the issue of making a valid inference with decision trees. A more conservative approach than BC is needed.

4.2 Proposed Method

4.2.1 Bias Correction on SD

Here we experimented again with the bootstrap correction approach in algorithm(3) using the three model configurations. With each given model, we generated training data \mathcal{D} of size $n = 500$ and test data set \mathcal{D}' of size $n' = 10,000$ with bootstrap sample $\mathcal{B}=500$.

Model A

Table 4.2: Table of result for biased corrected SD computed with training data \mathcal{D} and test data \mathcal{D}' .

Tree	node	n	\bar{y}_t	s_t	n'_t	\bar{y}'_t	s'_t	Bias	s''_t
1	4	71	2.9092	1.0897	1307	2.8963	1.3485	0.13048	1.2202
1	5	88	4.0702	1.38991	1834	4.3482	1.4219	0.17872	1.5686
1	8	42	4.7341	1.0273	734	4.9273	1.6052	0.17917	1.2064
1	9	42	6.3611	1.3996	759	6.4625	1.6082	0.2453	1.6449
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
100	34	10	10.7390	1.1267	172	9.7006	1.2046	0.2476	1.3744
100	37	18	9.2502	1.1510	420	9.3260	1.4322	0.3336	1.4847
100	38	12	11.2736	1.4251	243	10.8681	1.3455	0.3083	1.7334
100	39	19	12.2764	1.4389	416	11.3965	1.5548	0.3025	1.7414




Figure 4.2: Illustration of bias correction on SD s_t through model A. The green reference line is $y = x$.

Model B

Table 4.3: Table of result for biased corrected SD computed with training data \mathcal{D} and test data \mathcal{D}' .

Tree	node	n	\bar{y}_t	s_t	n'_t	\bar{y}'_t	s'_t	Bias	s''_t
1	4	47	5.3015	1.9158	821	5.6932	2.4249	0.3951	2.3109
1	6	17	6.1179	2.4292	424	6.6783	2.8131	0.4321	2.8613
1	7	12	12.0097	2.0447	290	10.9269	2.4160	0.4348	2.4795
1	10	15	5.8375	2.4362	345	5.6767	2.3768	0.3716	2.8078
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
100	24	32	14.5282	1.9469	451	14.7844	2.6364	0.3773	2.3243
100	27	22	14.4614	2.6309	361	14.4363	2.6073	0.3626	2.9936
100	28	86	16.8107	2.5779	1628	16.4838	2.7446	0.2833	2.8612
100	29	44	19.6521	2.6006	1011	19.1129	2.6582	0.3104	2.9111

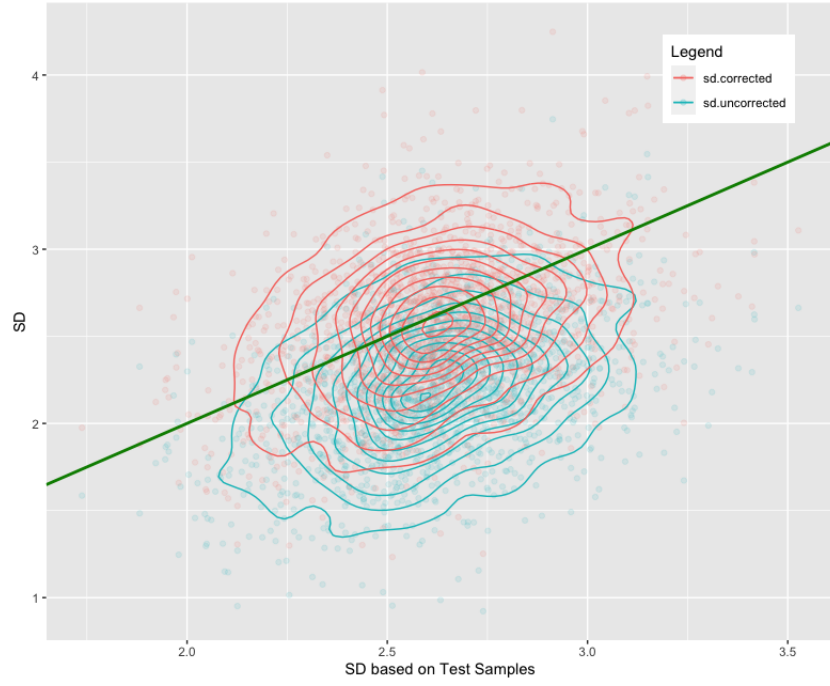


Figure 4.3: Illustration of bias correction on SD s_t through model B. The green reference line is $y = x$.

Model C

Table 4.4: Table of result for biased corrected SD computed with training data \mathcal{D} and test data \mathcal{D}' .

Tree	node	n	\bar{y}_t	s_t	n'_t	\bar{y}'_t	s'_t	Bias	s''_t
1	2	246	1.9491	0.9745	4690	2.0194	1.0145	0.04018	1.0147
1	4	126	1.8623	1.0429	2653	2.0107	1.0215	0.0410	1.0839
1	5	128	3.8070	0.9621	2657	3.8951	1.0985	0.0498	1.0119
2	2	228	1.9911	1.0627	4991	2.0207	0.9981	0.0394	1.1021
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
99	5	136	3.9981	1.0126	2543	3.9879	1.0232	0.0455	1.0581
100	2	267	2.0381	0.9767	5177	2.0212	1.0371	0.0400	1.0168
100	4	128	2.0268	1.0449	2424	1.9882	0.9991	0.0429	1.0878
100	5	105	3.8462	1.1190	2399	3.9983	0.9951	0.0589	1.1779




Figure 4.4: Illustration of bias correction on SD s_t through model C. The green reference line is $y = x$.

The best-sized tree \mathcal{T} is constructed with the simulated train data \mathcal{D} via pruning and cross-validation with the 1-SE rule (Breiman et al., 1984). The tree (Tree), the node, the total number of observations in each node (n) and the estimates $\{(n_t, \bar{y}_t, s_t) : t \in \tilde{\mathcal{T}}\}$ for all terminal nodes of \mathcal{T} are extracted and recorded in Tables 4.2, 4.3, and 4.4. Now on the basis of our simulated test data \mathcal{D}' , the estimates $\{(n'_t, \bar{y}'_t, s'_t) : t \in \tilde{\mathcal{T}}\}$ for all terminal nodes of \mathcal{T} are also extracted and recorded accordingly. Similarly, the bias estimates (Bias) and the bias-corrected SD (s''_t) in reference to Algorithm 3 our proposed method are as well recorded in each given table. Table 4.2, 4.3 and 4.4 present results from Model 4.1, 4.2 and 4.3 respectively.

Results from Table 4.2 and 4.3 show that s_t which is the naive standard error is lower

than that of s'_t obtained from the test sample. Hence using s_t directly to construct $(1 - \alpha)$ coverage for each terminal node will be over-optimistic. Specifically we estimated bias and added it to the s_t leading to the s''_t which corrects the downwards biasedness of the s_t . Also, Figures 4.2 and 4.3 are density contours that show the uncorrected SD s_t , corrected SD s''_t and the SD s'_t estimates from a large test sample \mathcal{D}' . We can see that the bias correction procedure really helps bring s_t up close to what they should be, namely, around s'_t computed from the test data. This is indicated by the line of reference from each plot. The reference line $y=x$ passes right through the center of density contours of (s'_t, s''_t) , but way above the density contours of (s'_t, s_t) . But model C which is the tree model has a perfect correspondence between the standard errors. Hence the simulated results show that our proposed methods work well in correcting the downwards biasedness of the s_t . The bias-correct SD estimates s''_t are more reliable for summarizing each terminal node.

4.2.2 Empirical Coverage

Now given the bootstrap bias corrected SD from the three models. We construct $(1 - \alpha) \times 100\%$ CI in node t for \bar{y}'_t based on s_t , s'_t and s''_t estimates with confidence level 95% to check whether or not the confidence bounds captures \bar{y}'_t estimate with appropriate coverages.

- (a) Naive: $\bar{y}_t \pm z_{1-\alpha/2}s_t/\sqrt{n_t}$;
- (b) BBC: $\bar{y}_t \pm z_{1-\alpha/2}s''_t/\sqrt{n_t}$;
- (c) Oracle: $\bar{y}_t \pm z_{1-\alpha/2}s'_t/\sqrt{n_t}$.

Model A

Table 4.5: Confident intervals using the s_t , s_t'' and s_t' from data generated by Model A

Tree	node	n	\bar{y}_t'	$\bar{y}_t \pm z_{1-\alpha/2} s_t / \sqrt{n_t}$	$\bar{y}_t \pm z_{1-\alpha/2} s_t'' / \sqrt{n_t}$	$\bar{y}_t \pm z_{1-\alpha/2} s_t' / \sqrt{n_t}$
1	4	71	2.896311	(2.655751, 3.162678)	(2.625402, 3.193028)	(2.595555, 3.222874)
1	5	88	4.348171	(3.779818, 4.360616)	(3.742478, 4.397956)	(3.773123, 4.367311)
1	8	42	4.927321	(4.423404, 5.044764)	(4.369219, 5.098949)	(4.24862, 5.219548)
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
100	37	18	9.325968	(8.718481, 9.781969)	(8.564353, 9.936098)	(8.588573, 9.911878)
100	38	12	10.86807	(10.467292, 12.07993)	(10.292874, 12.254349)	(10.512354, 12.034869)
100	39	19	11.39645	(11.629448, 12.923406)	(11.49341, 13.059443)	(11.577314, 12.975539)

Model B

Table 4.6: Confident intervals using the s_t , s_t'' and s_t' from data generated by Model B

Tree	node	n	\bar{y}_t'	$\bar{y}_t \pm z_{1-\alpha/2} s_t / \sqrt{n_t}$	$\bar{y}_t \pm z_{1-\alpha/2} s_t'' / \sqrt{n_t}$	$\bar{y}_t \pm z_{1-\alpha/2} s_t' / \sqrt{n_t}$
1	4	47	5.693237	(4.753805, 5.849232)	(4.640852, 5.962185)	(4.608278, 5.99476)
1	6	17	6.678329	(4.963116, 7.272624)	(4.757708, 7.478032)	(4.78061, 7.455129)
1	7	12	10.926954	(10.8528, 13.166542)	(10.606772, 13.412569)	(10.642701, 13.37664)
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
100	27	22	14.43631	(13.362, 15.56074)	(13.21046, 15.71228)	(13.37188, 15.55086)
100	28	86	16.48383	(16.26586, 17.35552)	(16.20599, 17.4154)	(16.23062, 17.39076)
100	29	44	19.11286	(18.88374, 20.42058)	(18.79201, 20.51231)	(18.86674, 20.43759)

Model C

Table 4.7: Confident intervals using the s_t , s_t'' and s_t' from data generated by Model C

Tree	node	n	\bar{y}_t'	$\bar{y}_t \pm z_{1-\alpha/2} s_t / \sqrt{n_t}$	$\bar{y}_t \pm z_{1-\alpha/2} s_t'' / \sqrt{n_t}$	$\bar{y}_t \pm z_{1-\alpha/2} s_t' / \sqrt{n_t}$
1	2	246	2.019431	(1.827346, 2.070895)	(1.822325, 2.075915)	(1.822342, 2.075898)
1	4	126	2.010694	(1.680181, 2.044396)	(1.67303, 2.051546)	(1.683922, 2.040655)
1	5	128	3.895063	(3.640335, 3.973686)	(3.631699, 3.982321)	(3.616707, 3.997313)
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
100	2	267	2.021208	(1.920895, 2.155209)	(1.916093, 2.160011)	(1.91366, 2.162444)
100	4	128	1.988171	(1.845818, 2.207855)	(1.838386, 2.215287)	(1.853749, 2.199924)
100	5	105	3.998295	(3.63215, 4.060223)	(3.620883, 4.07149)	(3.655842, 4.036531)

From Tables 4.5, 4.6 and 4.7 show 95% confidence intervals constructed with different SD estimates for each terminal of the final tree in each simulation run. A total of 100 simulation

runs are used.

4.2.3 Percentage Coverage Estimates

We send 1000 test samples of size 500 down each tree and check if the test sample means within each terminal node is covered by each 95% CI. The empirical coverage is essentially the relative frequency of when the 95 CI includes the test sample mean.

Table 4.8: Coverage Probabilities by BBC

Case	Naive (s_t)	BBC (s_t'')	Oracle(s_t')
Model A	0.8526	0.9117	0.9036
Model B	0.8509	0.9144	0.9283
Model C	0.9077	0.9240	0.9142

Values from Table 4.8 clearly indicate that the empirical coverage from the naive approach is far below the nominal level, i.e., 95%. Comparatively, the BBC approach, similar to the oracle approach, yields a coverage that is much closer.

Chapter 5

Real Data Analysis

5.1 Real Data Exploration

5.1.1 Data Source and Preparation

In this chapter, we apply our proposed method to real-life data as an illustration. For instance, as a statistical consultant or data scientist, you are tasked to provide statistical advice on how to estimate the salary of a particular player in a baseball team by prediction. The baseball team in question does not want to suggest a salary that is too high or too low. However, they know some of the characteristics of a previous team that influences their salary structure but would like an objective way of estimating the current and future salaries. The goal, therefore, is to develop accurate confidence bound estimates that can be used to determine a value within these bounds that can be used to predict a player's salary on the basis of the previous team's characteristics. Using data which provides information on Major League Baseball from the 1986 and 1987 seasons by Hitters, sourced from <http://lib.stat.cmu.edu/datasets/baseball.data> and also available from the ISLR package in R. The StatLib library at Carnegie Mellon University was the original host of this data set, which was also used in the 1988 ASA Graphics Section Poster Session. Essentially, this salary data was originally from Sports dated April 20, 1987, which captured excerpts on the 1986 and career statistics which were obtained from The 1987 Baseball Encyclopedia Update published by Collier Books, Macmillan Publishing Company, New York.(James et al., 2017).

The data contains the 1987 annual salary of baseball players (in thousands of dollars) on

the opening day of the season. It has 263 rows (observations) and 25 columns (variables). A brief description of the variables is provided in Table 5.1

Table 5.1: 1987 Baseball Salary Data for Hitters

Variable	Description
name	hitter's name
bat86	number of times at bat in 1986
hit86	number of hits in 1986
hr86	number of home runs in 1986
run86	number of runs in 1986
rb86	number of runs batted in in 1986
wlk86	number of walks in 1986
yrs	number of years in the major leagues
bater	number of times at bat during his career
hiter	number of hits during his career
hrer	number of home runs during his career
runcr	number of runs during his career
rber	number of runs batted in during his career
wlker	number of walks during his career
leag86	player's league at the end of 1986
div86	player's division at the end of 1986
team86	player's team at the end of 1986
pos86	player's position(s) in 1986
puto86	number of put outs in 1986
asst86	number of assists in 1986
err86	number of errors in 1986
salary	1987 annual salary on opening day in thousands of dollars
leag87	player's league at the beginning of 1987
team87	player's team at the beginning of 1987

Data preparation and validation were carried out. First, a salary which is the response

variable was transformed through log-transformation to make it less skewed. Again highly concentrated categorical and string independent variables such as name, team86, team 87, and pos86 were removed and less concentrated ones were recoded into 0's and 1's for further analysis.

5.1.2 Obtaining the best Tree from the data set

Final (1SE) Tree Model for 1987 Baseball Data




Figure 5.1: Plot of the best tree via pruning and cross validation.

Using the prepared data and the CART function. Tree analysis was performed through pruning and cross-validation. By the 1-SE rule (Breiman et al., 1984), the best tree was obtained as plotted in Fig(5.1).

5.1.3 Estimates of s_t and s_t^c from the best tree

Table 5.2: Table values of s_t and s_t'' using the baseball salary data.

node	n	\bar{y}_t	s_t	Bias	s_t''
3	39	4.548878	0.2439172	0.03598723	0.2799044
4	51	5.283612	0.3136554	0.06161825	0.3752736
6	112	6.256974	0.4964537	0.07495183	0.5714055
7	61	6.81953	0.4862067	0.08187276	0.5680795

Also, we applied the proposed BBC method with $B = 500$ bootstrap samples to estimate the biases and obtain the bias-corrected SD (s_t'') for each of the four terminal nodes. The results are tabulated in Table 5.3.

5.1.4 Empirical Coverage from the best tree

Table 5.3: Confidence estimates of the mean salary via BBC: $\bar{y}_t \pm z_{1-\alpha/2}s_t''/\sqrt{n_t}$

node	n	\bar{y}_t	$L_{\bar{y}_t}$	$U_{\bar{y}_t}$	$e^{\bar{y}_t}$	$L_{e^{\bar{y}_t}}$	$U_{e^{\bar{y}_t}}$
3	39	4.548878	4.461031	4.636724	94.52626	86.57672	103.2057
4	51	5.283612	5.180618	5.386605	197.08035	177.79261	218.4605
6	112	6.256974	6.151151	6.362798	521.6383	469.257	579.8667
7	61	6.81953	6.676972	6.962089	915.55493	793.91162	1055.8365

Table (5.3) show the 95% confidence interval for the node mean using our bias-corrected SD for all terminal nodes obtained from the best tree through our proposed method BBC: $\bar{y}_t \pm z_{1-\alpha/2}s_t''/\sqrt{n_t}$, where alpha was chosen to be $\alpha = 0.05$, n_t is the total number of samples in each individual terminal node, \bar{y}_t is logsalary (log of the mean salary) and s_t'' is the BBC SD estimates for each terminal node. For better interpretability, confidence intervals for

mean salary on its original scale are also obtained by taking the exponential. These are also shown in Table 5.3.




Figure 5.2: Plot of the Confident Estimates of the Mean Salary

Figure(5.2) is an error bound plot depicting table (5.3) estimates $(L_{e\tilde{y}_t}, U_{e\tilde{y}_t})$ graphically. Node 3 defines a group of players who have the lowest average salary. These players are characterized by the number of bats less than 1322 and the number of runs batted less than 56 in their career. Node 7 is characterized by players with the number of bats greater than 1322 and number of walks in 1986 greater than 53 in career having the highest average salary.

Chapter 6

Discussion and Conclusion

6.1 Significance of the Result

This research aimed to identify and correct the issue when a node-level inference is naively constructed with a $(1 - \alpha) \times 100\%$ confidence interval (CI) through Model (1.1). An interesting observation of downward biasedness is made concerning SD from the summary of the decision tree in Figure 3.1. Numerical and graphical results from our simulation study show that the BBC SD estimates are more reliable because they match well with the ‘gold’ SD estimates obtained from test data. Revealing that directly using SD (s_t) in Model (1.1) leads to an over-optimism of the CI estimates. However, we have been able to use a biased correction approach via bootstrapping to make correction. On this basis, a corrected SD (s_t^c) can now be used directly as indicated in Model (3.4). Hence, the bootstrap bias-corrected SD estimates become more convenient and applicable to use.

Another interesting observation is made in Table 4.8 which reveals that the bootstrap bias-corrected SD’s confidence interval estimate provides the highest coverage probability and that using the SD corrected obtained via bootstrapping yields a good confident bounds, this is further assessed in Table 5.3 of our real data analysis.

Therefore, making statistical inference on the confidence interval for the true node mean μ_t , from a decision tree model, which is the most common requests from users of decision trees, has been partly fulfilled in this research. We have focused on correction of the bias in the SD estimate. Constructing the confidence interval estimates from a tree model aids in predicting future values. However, one common issue is that constructing CI’s using relation 3.4 does not involve only the estimates from the summary of decision tree but also

a constant α and the choice of this α also plays an important role.

6.2 Recommendation

The primary aim of this project was to obtain more reliable and valid CI estimates for a tree model. As a result bootstrap bias correction approach was employed and the coverage probabilities were obtained. These coverage probabilities were obtained with regards to the coverage of each individual terminal node of the best tree summary. In effect, coverage across all terminal nodes was ignored. Even though our proposed method works better relative to the naïve method, the percentage coverage was not too convincing. We, therefore, recommend that, in future work, the multiplicity issues associated with the derivation of the CI estimates across all terminal nodes be addressed. Hence, we would like to explore the use of our proposed method combined with Scheffe or Tuckey's method within the one-way ANOVA setting or the bootstrap calibration (BC) method to handle the multiple comparisons across all terminal nodes in future research. Also, in future work, we recommend an extension of our approach to classification trees as well as the estimation of prediction intervals at each terminal node.

Appendix

R Codes

```
A # #####
# #####
# FUNCTION
# #####
# #####

library(rpart)

# =====
# GENERATE SOME DATA
# =====

rdat.MARS <- function(n, p=5, model="A")
{
  X <- NULL
  for (j in 1:p) {
    x <- runif(n)
    assign(paste("x", j, sep=""), x)
    X <- cbind(X, x)
  }
  if (model=="A") mu <- 0.1*exp(4*x1) + 4/(1+exp(-20*(x2-0.5))) + 3*x3 + 2*x4 + x5
  else if (model=="B") mu <- 10*sin(pi*x1*x2) + 20*(x3-0.5)^2 + 10*x4 + x5
  else if (model=="C") mu <- 2 + 2*sign(x1 <= 0.5)*sign(x2 <= 0.5)
  else stop("The arugment model= needs to be either A or B.")
  y <- mu + rnorm(n)
  dat <- data.frame(cbind(y, X))
  names(dat) <- c("y", paste("x", 1:p, sep=""))
  return(dat)
}

# =====
```

```

# FUNCTION cart() WRAPS UP STEPS FOR OBTAINING BEST TREE WITH rpart
# YET WITH FOCUS ON THE TERMINAL NODES ONLY
# =====

cart <- function(formula, data, method="anova", control=NULL,
size.selection=c("OSE", "1SE"), plot.it=FALSE){
  if (is.null(control)) control <- rpart.control(minsplit=20, minbucket=10,
maxdepth=5, cp=0, maxcompete=0, # NUMBER OF COMPETITIVE SPLITS
maxsurrogate=0, usesurrogate=2, surrogatestyle=0, # SURROGATE SPLITS FOR MISSING DATA
xval=10)
  tre0 <- rpart(formula=formula, data=data, method=method, control=control);
  if (size.selection=="OSE") {
    opt <- which.min(tre0$cptable[, "xerror"])
    best.cp <- tre0$cptable[opt, "CP"]; # print(cp.best)
    best.tree <- prune(tre0, cp = best.cp)
  } else if (size.selection=="1SE") {
    if (plot.it) plotcp(tre0, minline = TRUE) # 1SE
    cv.error <- (tre0$cptable)[,4]
    SE1 <- min(cv.error) + ((tre0$cptable)[,5])[which.min(cv.error)] # 1SE; CAN BE EASILY MODIFIED
    position <- min((1:length(cv.error))[cv.error <= SE1]); # print(position)
    # n.size <- (tre0$cptable)[,2] + 1 #TREE SIZE IS ONE PLUS NUMBER OF SPLITS.
    # best.size <- n.size[position]; # best.size
    best.cp <- sqrt(tre0$cptable[position,1]*tre0$cptable[(position-1),1]); # print(best.cp)
    # best.cp <- tre0$cptable[position,1]; print(best.cp)
    best.tree <- prune(tre0, cp=best.cp)
  }
  else stop("The values of size.selection= must be either OSE or 1SE")
  leaf.info <- best.tree$frame[best.tree$frame$var=="<leaf>", c(2, 4:5)]
  leaf.info$sd <- sqrt(leaf.info$dev/(leaf.info$n -1))
  # THE ROW NAMES DON'T MATCH WELL WITH TERMINAL NODES
  n.leaf <- aggregate(dat$y, by=list(best.tree$where), FUN=length); n.leaf
  leaf.info <- cbind(node=n.leaf$Group.1, leaf.info)
  # OUTPUT
  btree.size <- NROW(leaf.info)
  list(leaf=leaf.info, btree=best.tree, cp=best.cp, size=btree.size, tree0=tre0)
}

# =====
# SEND A TREE DOWN A DATASET AND RECOMPUTE

```

```
# =====

# LEAF INFO FROM fit.cart IS EXPANDED TO INCLUDE TEST SAMPLE INFO
send.down <- function(fit.cart, data, yname="y"){
  leaf <- fit.cart$leaf
  tree <- fit.cart$btrees
  node <- rpart:::pred.rpart(tree, x=rpart:::rpart.matrix(data));
  data$node <- node
  dat.tmp <- data[order(node), c(yname, "node")]
  leaf.test <- aggregate(dat.tmp$y, by=list(dat.tmp$node), FUN=length)
  yval.test <- aggregate(dat.tmp$y, by=list(dat.tmp$node), FUN=mean)$x
  sd.test <- aggregate(dat.tmp$y, by=list(dat.tmp$node), FUN=sd)$x
  # SUMMARIZE RESULTS
  leaf.test <- cbind(leaf.test, yval.test, sd.test)
  names(leaf.test) <- c("node", "n.test", "ybar.test", "sd.test")
  leaf.info <- merge(leaf, leaf.test, by="node", all.x = FALSE)
  return(leaf.info)
}
```

B

```
# #####
# TRIAL I: CHECK IF BOOTSTRAP CALIBRATION REALLY WORKS
# #####

rm(list=ls(all=TRUE))
source("R-FunctionsBC.R")

# set.seed(123)
nrun <- 3; B <- 500
n <- n.test <- 500; p <- 5; Model <- "A"
# n.test <- 2000;
alpha <- c(1:200/10000);

z0 <- qnorm(1-alpha/2); n.alpha <- length(alpha)
Alpha.True <- Alpha.Boots <- matrix(0, nrow=n.alpha, ncol=nrun)
for (i in 1:nrun){
  print(paste("===== run ", i, " =====", sep=""))
```

```

dat <- rdat.MARS(n=n, p=p, model=Model)
fit.tree <- cart(y~., data=dat, method="anova", size.selection="1SE");
best.tree <- fit.tree$btree
leaf <- fit.tree$leaf
P.True <- P.Boots <- matrix(0, nrow=n.alpha, ncol=B)

# POPULATION VERSION
for (b in 1:B){
test <- rdat.MARS(n=n.test, p=p, model=Model)
test.info <- senddown(tree=best.tree, data=test, yname="y")
for (k in 1:n.alpha){
LB <- leaf$yval - z0[k]*leaf$sd/sqrt(leaf$n)
UB <- leaf$yval + z0[k]*leaf$sd/sqrt(leaf$n)
lb <- ub <- factor(test.info$node, levels=leaf$node, ordered=TRUE)
levels(lb) <- LB; levels(ub) <- UB
lb <- as.numeric(as.character(lb)); ub <- as.numeric(as.character(ub))
P.True[k, b] <- mean((test.info$y >= lb) & (test.info$y <= ub))
}
}
# print(P.True)
Alpha.True[, i] <- apply(P.True, 1, mean)

# BOOTSTRAP CALIBRATION
for (b in 1:B){
id.b <- sample(1:n, size=n, replace=TRUE)
dat.b <- dat[id.b,]
fit.b <- cart(y~., data=dat.b, method="anova", size.selection="1SE");
btree.b <- fit.b$btree
leaf.b <- fit.b$leaf
info.b <- senddown(tree=btree.b, data=dat, yname="y")
for (k in 1:n.alpha){
LB <- leaf.b$yval - z0[k]*leaf.b$sd/sqrt(leaf.b$n)
UB <- leaf.b$yval + z0[k]*leaf.b$sd/sqrt(leaf.b$n)
lb <- ub <- factor(info.b$node, levels=leaf.b$node, ordered=TRUE)
levels(lb) <- LB; levels(ub) <- UB
lb <- as.numeric(as.character(lb)); ub <- as.numeric(as.character(ub))
P.Boots[k, b] <- mean((info.b$y >= lb) & (info.b$y <= ub))
}
}
# print(P.Boots)

```

```

Alpha.Boots[, i] <- apply(P.Boots, 1, mean)
}
apply(Alpha.True, 1, mean)
apply(Alpha.Boots, 1, mean)
save(Alpha.True, Alpha.Boots, file="Out-Trial-IA.Rdata")


# =====
# PLOTTING THE RESULTS
# =====

rm(list=ls(all=TRUE))
alpha <- c(1:50/10000);
postscript(file="fig-Trial-I.eps", horizontal=TRUE)
par(mfrow=c(1, 3), mar=c(8, 4, 8, 4))

# MODEL A
load("Out-trial-IA.Rdata")
M0 <- Alpha.Boots; dim(M0)
avg1.a <- apply(Alpha.Boots, 1, mean)
plot(x=range(alpha), y=c(.78, 1), type="n", xlab=expression(alpha),
ylab="Coverage Probability", main="Model A")
for (j in 1:NCOL(M0)){
a0 <- M0[,j]
lines(alpha, a0, lwd=0.005, col="orange")
}
M1 <- Alpha.True
avg2.a <- apply(M1, 1, mean)
for (j in 1:NCOL(M1)){
a0 <- M1[,j]
lines(alpha, a0, lwd=0.005, col="lightblue")
}
lines(alpha, avg1.a, lwd=1.5, col="brown")
lines(alpha, avg2.a, lwd=1.5, col="blue")
abline(h=0.95, col="gray50", lwd=0.8, lty=2)
legend(0.001, 1.00, legend=c("Bootstrap", "Population"), lty=1,

```

```

col=c("orange", "blue"), lwd=1)

# MODEL B
load("Out-trial-IB.Rdata")
M0 <- Alpha.Boots; dim(M0)
avg1.a <- apply(Alpha.Boots, 1, mean)
plot(x=range(alpha), y=c(.76, 1), type="n", xlab=expression(alpha),
ylab="Coverage Probability", main="Model B")
for (j in 1:NCOL(M0)){
  a0 <- M0[,j]
  lines(alpha, a0, lwd=0.005, col="orange")
}
M0 <- Alpha.True
avg2.a <- apply(M0, 1, mean)
for (j in 1:NCOL(M0)){
  a0 <- M0[,j]
  lines(alpha, a0, lwd=0.005, col="lightblue")
}
lines(alpha, avg1.a, lwd=1.5, col="brown")
lines(alpha, avg2.a, lwd=1.5, col="blue")
abline(h=0.95, col="gray50", lwd=0.8, lty=2)

# MODEL C
load("Out-trial-IC.Rdata")
M0 <- Alpha.Boots; dim(M0)
dat.tmp <- data.frame(cbind(alpha, M0))

library(tidyverse)
dat.tmp

avg1.a <- apply(Alpha.Boots, 1, mean)
plot(x=range(alpha), y=c(.74, 1), type="n", xlab=expression(alpha),
ylab="Coverage Probability", main="Model C")
for (j in 1:NCOL(M0)){
  a0 <- M0[,j]
  lines(alpha, a0, lwd=0.005, col="orange")
}
M0 <- Alpha.True

```



```

avg2.a <- apply(M0, 1, mean)
for (j in 1:NCOL(M0)){
  a0 <- M0[,j]
  lines(alpha, a0, lwd=0.005, col="lightblue")
}
lines(alpha, avg1.a, lwd=1.5, col="brown")
lines(alpha, avg2.a, lwd=1.5, col="blue")
abline(h=0.95, col="gray50", lwd=0.8, lty=2)

dev.off()

```

```

C #####
# BIAS-CORRECTION OF SD IN DECISION TREES
#####

# -----
# MODEL A
# -----

rm(list=ls(all=TRUE))
source("Functions-BBC.R")

set.seed(111)
nrun <- 100
B <- 200; n <- 500; Model <- "A";
# positive.bias <- FALSE
positive.bias <- TRUE
p <- 5; n0 <- 10000;
OUT <- NULL
TREE <- as.list(1:nrun)
for (i in 1:nrun) {
  dat <- rdat.MARS(n=n, p=p, model=Model)
  fit.cart <- cart(y~., data=dat, method="anova", size.selection="1SE", plot.it=FALSE);
  TREE[[i]] <- fit.cart$btrees
  node.0 <- rpart::pred.rpart(fit.cart$btrees, x=rpart::rpart.matrix(dat));
  test <- rdat.MARS(n=n0, p=p, model=Model)
  info.0 <- send.down(fit.cart, data=test, yname="y");
  sd.un <- info.0$sd

```

```

# BOOTSTRAP CORRECTION
bias <- rep(0, length(sd.un))
for (b in 1:B){
  print(cbind(run=i, boots=b))
  id.b <- sample(1:n, size=n, replace=TRUE)
  dat.b <- dat[id.b,]; dat.oob <- dat[-unique(id.b),]
  fit.b <- cart(y~., data=dat.b, method="anova", size.selection="1SE", plot.it=FALSE);
  info.b <- send.down(fit.b, data=dat, yname="y") ## SHOULD USE dat.oob?
  bias.b <- info.b$sd.test - info.b$sd
  if (positive.bias) bias.b <- pmax(bias.b, 0) ### NECESSARY?
  node.b <- rpart::pred.rpart(fit.b$btree, x=rpart::rpart.matrix(dat));
  tab <- table(node.0, node.b)
  M.prop <- prop.table(tab, 1)
  bias.b <- M.prop%*%bias.b
  bias <- bias + bias.b
}
bias <- bias/B
sd.co <- sd.un + bias
out <- cbind(tree=i, info.0,bias, sd.co)
OUT <- rbind(OUT, out)
}
OUT <- as.data.frame(OUT)
colnames(OUT) <- c("tree", "node", "n", "dev", "ybar", "sd.uncorrected",
  "n.test", "ybar.test", "sd.test",
  "bias", "sd.corrected")
head(OUT)
save(OUT, TREE, file="result-ModelA.Rdat")

# -----
# MODEL B
# -----

rm(list=ls(all=TRUE))
source("Functions-BBC.R")

set.seed(111)
nrun <- 100
B <- 200; n <- 500; Model <- "B";

```

```

# positive.bias <- FALSE
positive.bias <- TRUE
p <- 5; n0 <- 10000;
OUT <- NULL
TREE <- as.list(1:nrun)
for (i in 1:nrun) {
  dat <- rdat.MARS(n=n, p=p, model=Model)
  fit.cart <- cart(y~., data=dat, method="anova", size.selection="1SE", plot.it=FALSE);
  TREE[[i]] <- fit.cart$btrees
  node.0 <- rpart::pred.rpart(fit.cart$btrees, x=rpart::rpart.matrix(dat));
  test <- rdat.MARS(n=n0, p=p, model=Model)
  info.0 <- send.down(fit.cart, data=test, yname="y");
  sd.un <- info.0$sd

  # BOOTSTRAP CORRECTION
  bias <- rep(0, length(sd.un))
  for (b in 1:B){
    print(cbind(run=i, boots=b))
    id.b <- sample(1:n, size=n, replace=TRUE)
    dat.b <- dat[id.b,]; dat.oob <- dat[-unique(id.b),]
    fit.b <- cart(y~., data=dat.b, method="anova", size.selection="1SE", plot.it=FALSE);
    info.b <- send.down(fit.b, data=dat, yname="y") ## SHOULD USE dat.oob?
    bias.b <- info.b$sd.test - info.b$sd
    if (positive.bias) bias.b <- pmax(bias.b, 0) ### NECESSARY?
    node.b <- rpart::pred.rpart(fit.b$btrees, x=rpart::rpart.matrix(dat));
    tab <- table(node.0, node.b)
    M.prop <- prop.table(tab, 1)
    bias.b <- M.prop%*%bias.b
    bias <- bias + bias.b
  }
  bias <- bias/B
  sd.co <- sd.un + bias
  out <- cbind(tree=i, info.0, bias, sd.co)
  OUT <- rbind(OUT, out)
}
OUT <- as.data.frame(OUT)
colnames(OUT) <- c("tree", "node", "n", "dev", "ybar", "sd.uncorrected",
  "n.test", "ybar.test", "sd.test",
  "bias", "sd.corrected")
head(OUT)

```

```

save(OUT, TREE, file="result-ModelB.Rdat")

# -----
# MODEL C (TRUE TREE)
# -----

rm(list=ls(all=TRUE))
source("Functions-BBC.R")

set.seed(111)
nrun <- 100
B <- 200; n <- 500; Model <- "C";
# positive.bias <- FALSE
positive.bias <- TRUE
p <- 5; n0 <- 10000;
OUT <- NULL
TREE <- as.list(1:nrun)
for (i in 1:nrun) {
  dat <- rdat.MARS(n=n, p=p, model=Model)
  fit.cart <- cart(y~., data=dat, method="anova", size.selection="1SE", plot.it=FALSE);
  TREE[[i]] <- fit.cart$btrees
  node.0 <- rpart::pred.rpart(fit.cart$btrees, x=rpart::rpart.matrix(dat));
  test <- rdat.MARS(n=n0, p=p, model=Model)
  info.0 <- send.down(fit.cart, data=test, yname="y");
  sd.un <- info.0$sd

  # BOOTSTRAP CORRECTION
  bias <- rep(0, length(sd.un))
  for (b in 1:B){
    print(cbind(run=i, boots=b))
    id.b <- sample(1:n, size=n, replace=TRUE)
    dat.b <- dat[id.b,]; dat.oob <- dat[-unique(id.b),]
    fit.b <- cart(y~., data=dat.b, method="anova", size.selection="1SE", plot.it=FALSE);
    info.b <- send.down(fit.b, data=dat, yname="y") ## SHOULD USE dat.oob?
    bias.b <- info.b$sd.test - info.b$sd
    if (positive.bias) bias.b <- pmax(bias.b, 0) ### NECESSARY?
    node.b <- rpart::pred.rpart(fit.b$btrees, x=rpart::rpart.matrix(dat));
    tab <- table(node.0, node.b)
    M.prop <- prop.table(tab, 1)
  }
}

```

```

bias.b <- M.prop%*%bias.b
bias <- bias + bias.b
}
bias <- bias/B
sd.co <- sd.un + bias
out <- cbind(tree=i, info.0,bias, sd.co)
OUT <- rbind(OUT, out)
}
OUT <- as.data.frame(OUT)
colnames(OUT) <- c("tree", "node", "n", "dev", "ybar", "sd.uncorrected",
"n.test", "ybar.test", "sd.test",
"bias", "sd.corrected")
head(OUT)
save(OUT, TREE, file="result-ModelC.Rdat")

#####
# EXPLORING THE REULSTS
#####

rm(list=ls(all=TRUE))
library(tidyverse)

load("result-ModelA.Rdat")
#load("result-ModelA-0SE.Rdat")
#load("result-ModelB.Rdat")
#load("result-ModelC.Rdat")

ls()
names(OUT); head(OUT)
tail(OUT)

OUT %>%
select(tree, node, sd.uncorrected, sd.test, sd.corrected) %>%
ggplot() +
geom_point(aes(x=sd.test, y=sd.uncorrected, colour="sd.uncorrected"), alpha=0.1) +
#xlab("SD based on Test Samples") + ylab("SD") +
geom_point(aes(x=sd.test, y=sd.corrected, colour="sd.corrected"), alpha=0.1) +
geom_density_2d(aes(x=sd.test, y=sd.uncorrected, colour="sd.uncorrected") ) +

```

```

geom_density_2d(aes(x=sd.test, y=sd.corrected, color="sd.corrected") ) +
geom_abline(slope=1, intercept = 0, colour="green4", size=1)+
labs(
  x = "SD based on Test Samples",
  y = "SD",
  #title = "Plot of the Uncorrected and Bias-Corrected SD",
  color = "Legend")+ theme(
  legend.position = c(0.95, 0.95),
  legend.justification = c("right", "top"),
  legend.box.just = "right",
  legend.margin = margin(6, 6, 6, 6)
)

```

```

D # =====
# COVERAGE
# =====

```

```

rm(list=ls(all=TRUE))
#source("Functions-BBC.R")

```

```

#load("result-ModelA.Rdat")
#load("result-ModelB.Rdat")
load("result-ModelC.Rdat")
TREE <- BTREE

```

```

conf.level <- 0.95
alpha <- (1-conf.level)
OUT %>%
#na.exclude() %>%
mutate(L.naive = ybar - qnorm(1-alpha/2) * sd.uncorrected/sqrt(n),
  U.naive = ybar + qnorm(1-alpha/2) * sd.uncorrected/sqrt(n),
  L.BBC = ybar - qnorm(1-alpha/2) * sd.corrected/sqrt(n),
  U.BBC = ybar + qnorm(1-alpha/2) * sd.corrected/sqrt(n),
  L.oracle = ybar - qnorm(1-alpha/2) * sd.test/sqrt(n),
  U.oracle = ybar + qnorm(1-alpha/2) * sd.test/sqrt(n)) %>%
select(tree, node, n, L.naive, U.naive, L.BBC, U.BBC, L.oracle, U.oracle) -> CI

CI %>% tail()

```

```

n.trees <- 10
n.sample <- 1000
n0 <- 10000; p <- 5
Model <- "C"
yname <- "y"
COVER <- NULL
for (i in 1:n.trees){
  tree.i <- TREE[[i]]
  CI.i <- CI %>% filter(tree==i)
  cover.naive <- cover.BBC <- cover.oracle <- rep(0, NROW(CI.i))
  for (j in 1:n.sample) {
    print(cbind(tree=i, sample=j))
    dat <- rdat.MARS(n=n0, p=p, model=Model)
    node <- rpart::pred.rpart(tree.i, x=rpart::rpart.matrix(dat));
    dat$node <- node
    dat.tmp <- dat[order(node), c(yname, "node")]
    ybar.test <- aggregate(dat.tmp$y, by=list(dat.tmp$node), FUN=mean)$x
    cover.naive <- cover.naive + sign(ybar.test >= CI.i$L.naive & ybar.test <= CI.i$U.naive)
    cover.BBC <- cover.BBC + sign((ybar.test >= CI.i$L.BBC) & (ybar.test <= CI.i$U.BBC))
    cover.oracle <- cover.oracle + sign(ybar.test >= CI.i$L.oracle & ybar.test <= CI.i$U.oracle)
  }
  CI.i %>% mutate(cover.naive=cover.naive/n.sample,
    cover.BBC=cover.BBC/n.sample,
    cover.oracle=cover.oracle/n.sample) -> CI.i
  COVER <- rbind(COVER, CI.i)
}
apply(COVER, 2, mean)

```

E

```
##### Real Data Exploration #####
```

```

rm(list=ls(all=TRUE))
#setwd("~/Desktop/THESIS U/Updated/real data")
source("Functions-BBC.R")
baseball <- read.table("bb87.dat", header = F,
  col.names=c("id", "name", "bat86", "hit86", "hr86", "run86", "rb86", "wlk86",
    "yrs", "batcr", "hitcr", "hrcr", "runcr", "rbcr", "wlkcr", "leag86", "div86",

```

```

"team86", "pos86", "puto86", "asst86", "err86","salary", "leag87", "team87",
"logsalary"))

apply(baseball, 2, FUN=function(x) length(unique(x)))

dat <- baseball %>%
mutate(y=logsalary, team.change=sign(team86!=team87),
leag86=sign(leag86=="A"),
leag87=sign(leag87=="A"),
div86=sign(div86=="W")) %>%
select(-salary, -id, -name, -logsalary, -team86, -team87, -pos86) %>%
select(y, everything()) %>%
as.data.frame()
head(dat)
anyNA(dat)

#####
fit.cart <- cart(y ~., data=dat, method="anova",
size.selection="1SE", plot.it=TRUE,model= TRUE);
btree <- fit.cart$btree
node.0 <- rpart::pred.rpart(btree, x=rpart::rpart.matrix(dat));
info.0 <- fit.cart$leaf
sd.un <- info.0$sd

#### Exploring the btree via a plot #####
library(rpart.plot)
library(RColorBrewer)
rpart.plot(btree, shadow.col="gray", extra=1,
main="Final (1SE) Tree Model for 1987 Baseball Data")

# BOOTSTRAP CORRECTION
B <- 500
n <- nrow(dat)
positive.bias <- TRUE
bias <- rep(0, length(sd.un))
for (b in 1:B){
print(paste("=====", b, "=====", sep=""))
id.b <- sample(1:n, size=n, replace=TRUE)

```



```

dat.b <- dat[id.b,]
fit.b <- cart(y ~., data=dat.b, method="anova",
size.selection="1SE", plot.it=FALSE);
info.b <- send.down(fit.b, data=dat, yname="y")
bias.b <- info.b$sd.test - info.b$sd
if (positive.bias) bias.b <- pmax(bias.b, 0)
node.b <- rpart:::pred.rpart(fit.b$btrees, x=rpart:::rpart.matrix(dat));
tab <- table(node.0, node.b)
M.prop <- prop.table(tab, 1)
bias.b <- M.prop%*%bias.b
bias <- bias + bias.b
}
bias <- bias/B
sd.co <- sd.un + bias
out <- cbind(info.0,bias, sd.co)
OUT <- as.data.frame(out)
colnames(OUT) <- c("node", "n", "dev", "ybar", "sd.uncorrected","bias", "sd.corrected")
save(OUT, file="result-bb.Rdat")

# =====
#   COVERAGE           #
# =====

# load("result-bb.Rdat")
conf.level <- 0.95
alpha <- (1-conf.level)
OUT %>%
mutate(L.naive = ybar - qnorm(1-alpha/2) * sd.uncorrected/sqrt(n),
U.naive = ybar + qnorm(1-alpha/2) * sd.uncorrected/sqrt(n),
L.BBC = ybar - qnorm(1-alpha/2) * sd.corrected/sqrt(n),
U.BBC = ybar + qnorm(1-alpha/2) * sd.corrected/sqrt(n),
ybar.sal = exp(ybar),
L.sal = exp(L.BBC),
U.sal=exp(U.BBC))%>%
select(node, n, ybar, L.BBC, U.BBC, ybar.sal, L.sal, U.sal) -> CI
CI
save(CI, file="CI-bb.Rdat")

```

```

CI %>%
mutate(node=factor(node))
ggplot(CI, aes(x=node, y=ybar.sal, group = node)) +
geom_errorbar(aes(ymin=L.sal, ymax=U.sal), color="blue") +
geom_point(size=5, color="tomato")

```

```

pd <- position_dodge(0.70)
ggplot(CI, aes(x=node, y = ybar.sal, group = node)) +
geom_point(position=pd) +
geom_errorbar(data=CI, aes(ymin=L.sal, ymax=U.sal,
color=node), width=1, position=pd)

```

References

- Breiman, L. (1996). Bagging predictors. *Machine learning*, 24(2):123–140.
- Breiman, L. (2001). Random forests. *Machine learning*, 45(1):5–32.
- Breiman, L., Friedman, J., Stone, C. J., and Olshen, R. A. (1984). *Classification and regression trees*. CRC press.
- de Ville, B. (2013). Decision trees. *Wiley Interdisciplinary Reviews: Computational Statistics*, 5(6):448–455.
- Efron, B. (1979). The 1977 rietz lecture. *The Annals of Statistics*, 7(1):1–26.
- Freund, Y., Schapire, R. E., et al. (1996). Experiments with a new boosting algorithm. In *icml*, volume 96, pages 148–156. Citeseer.
- Friedman, J. H. (1991). Multivariate adaptive regression splines. *The annals of statistics*, pages 1–67.
- Hu, X., Rudin, C., and Seltzer, M. (2019). Optimal sparse decision trees. *Advances in Neural Information Processing Systems (NeurIPS)*.
- James, G., Witten, D., Hastie, T., and Tibshirani, R. (2017). Data for an introduction to statistical learning with applications in r. *Package ‘ISLR.’CRAN*.
- Jiao, J. and Han, Y. (2017). Bias correction with jackknife, bootstrap, and taylor series. *arXiv e-prints*, pages arXiv–1709.
- Jordan, M. I. and Jacobs, R. A. (1994). Hierarchical mixtures of experts and the em algorithm. *Neural computation*, 6(2):181–214.
- Lakshminarayanan, B. (2016). *Decision trees and forests: a probabilistic perspective*. PhD thesis, UCL (University College London).
- Loh, W., Man, M., and Wang, S. (2018). Subgroups from regression trees with adjustment for prognostic effects and postselection inference. *Statistics in Medicine*, 38(4):545–557.
- Loh, W.-Y. (1987). Calibrating confidence coefficients. *Journal of the American Statistical Association*, 82(397):155–162.
- Loh, W.-Y. (1991). Bootstrap calibration for confidence interval construction and selection. *Statistica Sinica*, pages 477–491.
- Maimon, O. Z. and Rokach, L. (2014). *Data mining with decision trees: theory and applications*, volume 81. World scientific.

- Murthy, K. V. S. and Salzberg, S. L. (1995). *On growing better decision trees from data*. PhD thesis, Citeseer.
- Njoku, O. C. (2019). Decision trees and their application for classification and regression problems.
- Quinlan, J. (1993). Program for machine learning. *C4*. 5.
- Rokach, L. and Maimon, O. (2005). Decision trees. In *Data mining and knowledge discovery handbook*, pages 165–192. Springer.
- Su, X., Tsai, C.-L., Wang, H., Nickerson, D. M., and Li, B. (2009). Subgroup analysis via recursive partitioning. *Journal of Machine Learning Research*, 10(2).
- Su, X., Zhou, T., Yan, X., Fan, J., and Yang, S. (2008). Interaction trees with censored survival data. *The international journal of biostatistics*, 4(1).
- Tibshirani, R. J. and Efron, B. (1993). An introduction to the bootstrap. *Monographs on statistics and applied probability*, 57:1–436.
- Xia, L.-y., Zakowski, Y., He, P., Hur, C.-K., Malecha, G., Pierce, B. C., and Zdancewic, S. (2019). Interaction trees: representing recursive and impure programs in coq. *Proceedings of the ACM on Programming Languages*, 4(POPL):1–32.

Vita

George Ekow Quaye was born on May 16, 1991. He graduated from Adidome Senior High School, Ghana in 2011. George was a beneficiary, Government Scholarship for Best Science Student, and also earn an award for the best graduating science student in Adidome Senior High School. George, having taken an equal number of classes in Statistics, Economics, and Mathematics as an Actuarial Science student and excelled in them, successfully graduated in 2016 with a Bachelor of Science (BS) degree in Actuarial Science at the University of Cape Coast. This feat, coupled with his dutiful nature and research interests, earned him a one-year position as a Research Assistant in one of the school's prestigious departments, The Department of Clinical Nutrition and Dietetics. With his background, he tutored in courses related to statistics, he collected, analyzed, and interpreted statistical data for four different projects, two of which are: *The Prevalence of Anaemia in Adolescent Girls and Women in Ghana* with Michigan State University in collaboration with Ghana's flagship the University of Ghana as well as, *Iron Status and Psychosocial Well-being of Pregnant Women in Ghana* with Pennsylvania State University in collaboration with the University of Cape Coast. These projects helped him to practically acquire the systematic procedure, knowledge, and technical skills in collecting both qualitative and quantitative data, organizing, analyzing using the appropriate statistical tool, and giving interpretation based on the output, and making necessary inferences. Also, as a supervisor on one of these taught him how to lead, plan, schedule, and interviewed a participant in data collection. He is currently a graduate student majoring in Statistics at the University of Texas at El Paso (UTEP) and working at the Department of Mathematical Sciences of UTEP as a graduate teaching assistant, where he tutor at the Mathematics Resource Center for Students (MARCS) and grade scripts for assigned Professors and classes and expected to graduate in May 2021. As required by a graduate study, he started working on a research topic, *Making valid inference with decision trees* as his thesis with one of the elite professors in the department in person of Professor Xiaogang Su in his second year of Master's Studies. George will pursue his doctoral degree in Data Science at the University of Texas at El Paso (UTEP) in El Paso where he is currently obtaining his Master's degree in Statistics. Email: gequaye@miners.utep.edu